

```
to.encoder=False, enable_pdmux=False, pdmux_config_path=None, sm_group_num=8, custom_weight_loader=[], weight_loader_disable_m
weight_loa
vice_port
ine=False
draft_co
rocess_co
p_config=
ing defau
to-detect
rser=qwen
it torch
it torch
ad weight
rds: 0%
rds: 100%
ad weight
Cache is
mory pool
pture tar
52, 384,
m_tokens=
biling a
mpiling a
m_tokens=
m_tokens=
pture tar
pture tar
avail_mem
pture tar
x_total_num_tokens=20909, chunked_prefill_size=2048, max_prefill_tokens=16384, max_running_requests=2048, context_len=40960, a
ee cache initialized: source=default impl=RadixCache hybrid_swa=False hybrid_ssm=False hierarchical=False streaming_wrapped=Fa
F0: Started server process [5203]
F0: Waiting for application startup.
```

```
linlin@LAPTOP-DJ6RM1: ~$ curl -X POST http://localhost:30000/v1/chat/completions \
-H "Content-Type: application/json" \
-d '{
  "model": "Qwen/Qwen3-0.6B",
  "messages": [
    {"role": "user", "content": "你好，请简单介绍一下你自己。"}
  ],
  "max_tokens": 128
}'

[[{"id": "bd5ed55e3a9c4e56a05bece1e6bee520", "object": "chat.completion", "created": 1790317498, "model": "Qwen/Qwen3-0.6B", "
choices": [{"index": 0, "message": {"role": "assistant", "content": "\n\n您好！我是大模型，可以协助您完成各种任务，比如回答问题
、创作内容等。需要我帮您做什么呢？", "reasoning_content": "\n\n好的，用户问我的简单介绍。我需要先确认用户需要的是我的基本信
息，比如名字、身份，还是其他方面的信息。首先，我应该礼貌地回答，说明我是大模型，然后简要介绍我的功能和特点。要注意保持回
答简洁，避免冗长。用户可能想了解我的能力范围，所以需要明确说明我可以帮助的问题类型。同时，要确保回答友好且专业，让用户觉
得有帮助。"}, {"role": "assistant", "content": "\n\n好的，用户问我的简单介绍。我需要先确认用户需要的是我的基本信
息，比如名字、身份，还是其他方面的信息。首先，我应该礼貌地回答，说明我是大模型，然后简要介绍我的功能和特点。要注意保持回
答简洁，避免冗长。用户可能想了解我的能力范围，所以需要明确说明我可以帮助的问题类型。同时，要确保回答友好且专业，让用户觉
得有帮助。"}], "total_tokens": 137, "completion_tokens": 123, "prompt_tokens_details": null, "reasoning_tokens": 93, "metadata": {"weight_v
ersion": "default"}}]
linlin@LAPTOP-DJ6RM1: ~$
linlin@LAPTOP-DJ6RM1: ~$ ~

(Cray) linlin@LAPTOP-DJ6RM1: ~/test$ python build_workload2.py
Request 0 completed. TTFT: 1.4597s, Total Latency: 1.7248s
Request 1 completed. TTFT: 2.9759s, Total Latency: 3.2710s
Request 10 failed: Error code: 400 - {'object': 'error', 'message': 'The input (62747 tokens) is longer than the model's context length (40960 tokens).', 'type': 'BadRequestError', 'param': None, 'code': 400}
Request 18 failed: Error code: 400 - {'object': 'error', 'message': 'The input (69142 tokens) is longer than the model's context length (40960 tokens).', 'type': 'BadRequestError', 'param': None, 'code': 400}
Request 3 completed. TTFT: 2.8230s, Total Latency: 7.1431s
Request 2 completed. TTFT: 2.7771s, Total Latency: 9.6448s
Request 5 completed. TTFT: 8.3163s, Total Latency: 8.6143s
Request 4 completed. TTFT: 7.6254s, Total Latency: 10.7880s
Request 7 completed. TTFT: 11.4813s, Total Latency: 13.9715s
Request 8 completed. TTFT: 15.5963s, Total Latency: 15.6212s
Request 6 completed. TTFT: 8.0811s, Total Latency: 16.1315s
Request 9 completed. TTFT: 16.3107s, Total Latency: 16.5881s
Request 13 failed: Input length (37813 tokens) exceeds the maximum allowed length (20903 tokens). Use a shorter input or enable --allow-auto-truncate.
Request 11 completed. TTFT: 17.0616s, Total Latency: 17.5926s
Request 15 failed: Input length (23496 tokens) exceeds the maximum allowed length (20903 tokens). Use a shorter input or enable --allow-auto-truncate.
Request 12 completed. TTFT: 16.7675s, Total Latency: 17.0923s
Request 16 completed. TTFT: 17.7858s, Total Latency: 18.3168s
Request 19 failed: Input length (38171 tokens) exceeds the maximum allowed length (20903 tokens). Use a shorter input or enable --allow-auto-truncate.
Request 17 completed. TTFT: 19.3311s, Total Latency: 19.7735s
Request 14 completed. TTFT: 16.6648s, Total Latency: 22.0062s

=== Workload Results ===
request_id input_tokens output_tokens status TTFT(s) Latency(s)
0 6169 27 success 1.4597 1.7248
1 6237 3 success 2.9759 3.2710
10 20617 438 failed: Error code: 400 - {'object': 'error', 'message': 'The input (62747 tokens) is longer than the model's context length (40960 tokens).', 'type': 'BadRequestError', 'param': None, 'code': 400} NaN NaN
18 32615 392 failed: Error code: 400 - {'object': 'error', 'message': 'The input (69142 tokens) is longer than the model's context length (40960 tokens).', 'type': 'BadRequestError', 'param': None, 'code': 400} NaN NaN
3 6749 547 success 2.8230 7.1431
2 2445 521 success 2.7771 9.6448
5 3145 23 success 8.3163 8.6143
4 5002 231 success 7.6254 10.7880
7 6490 207 success 11.4813 13.9715
8 7280 3 success 15.5963 15.6212
6 910 515 success 8.0811 16.1315
9 6170 27 success 16.3107 16.5881
13 17493 648 failed: Input length (37813 tokens) exceeds the maximum allowed length (20903 tokens). Use a shorter input or enable --allow-auto-truncate. NaN NaN
11 6368 3 success 17.0616 17.5926
15 11132 422 failed: Input length (23496 tokens) exceeds the maximum allowed length (20903 tokens). Use a shorter input or enable --allow-auto-truncate. NaN NaN
12 3119 17 success 16.7675 17.0923
16 7174 42 success 17.7858 18.3168
19 18054 269 failed: Input length (38171 tokens) exceeds the maximum allowed length (20903 tokens). Use a shorter input or enable --allow-auto-truncate. NaN NaN
17 6423 37 success 19.3311 19.7735
14 1343 273
```